

Mod13_Sesion03MLibArbolClasificacion

September 23, 2023

1 Árboles de Clasificación

```
[1]: from pyspark import SparkContext
      from pyspark.sql import SQLContext
```

```
[6]: #Leer el contenido de una carpeta
      #Para leer de HDFS usar
      # hdfs:///tmp/dcd/OnTimeDB
      #Para leer de local usar
      # file:/home/cloudera/dcd/OnTimeDB/

      ## Descargar archivo zip y subir al cluster
      !wget https://github.com/omarmendoza564/datos/raw/main/datos/OnTimeDB.zip -O /
      ↪home/sergio_ibarra1795/OnTimeDB.zip
      !unzip -o /home/sergio_ibarra1795/OnTimeDB.zip -d /home/sergio_ibarra1795/
      !ls -la /home/sergio_ibarra1795/OnTimeDB
      !hdfs dfs -mkdir /tmp/dcd/OnTimeDB
      !hdfs dfs -put /home/sergio_ibarra1795/OnTimeDB/ /tmp/dcd/
```

--2023-09-23 18:24:40--

https://github.com/omarmendoza564/datos/raw/main/datos/OnTimeDB.zip

Resolving github.com (github.com)... 140.82.112.3

Connecting to github.com (github.com)|140.82.112.3|:443... connected.

HTTP request sent, awaiting response... 302 Found

Location:

https://raw.githubusercontent.com/omarmendoza564/datos/main/datos/OnTimeDB.zip

[following]

--2023-09-23 18:24:40--

https://raw.githubusercontent.com/omarmendoza564/datos/main/datos/OnTimeDB.zip

Resolving raw.githubusercontent.com (raw.githubusercontent.com)...

185.199.111.133, 185.199.108.133, 185.199.109.133, ...

Connecting to raw.githubusercontent.com

(raw.githubusercontent.com)|185.199.111.133|:443... connected.

HTTP request sent, awaiting response... 200 OK

Length: 484951 (474K) [application/zip]

Saving to: '/home/sergio_ibarra1795/OnTimeDB.zip'

/home/sergio_ibarra 100%[=====>] 473.58K --.-KB/s in 0.03s

2023-09-23 18:24:40 (13.3 MB/s) - '/home/sergio_ibarra1795/OnTimeDB.zip' saved
[484951/484951]

Archive: /home/sergio_ibarra1795/OnTimeDB.zip

```
  creating: /home/sergio_ibarra1795/OnTimeDB/
  inflating: /home/sergio_ibarra1795/OnTimeDB/.part-00000.crc
  inflating: /home/sergio_ibarra1795/OnTimeDB/.part-00001.crc
  inflating: /home/sergio_ibarra1795/OnTimeDB/.part-00002.crc
  inflating: /home/sergio_ibarra1795/OnTimeDB/.part-00003.crc
  inflating: /home/sergio_ibarra1795/OnTimeDB/.part-00004.crc
  inflating: /home/sergio_ibarra1795/OnTimeDB/.part-00005.crc
  inflating: /home/sergio_ibarra1795/OnTimeDB/.part-00006.crc
  inflating: /home/sergio_ibarra1795/OnTimeDB/part-00000
  inflating: /home/sergio_ibarra1795/OnTimeDB/part-00001
  inflating: /home/sergio_ibarra1795/OnTimeDB/part-00002
  inflating: /home/sergio_ibarra1795/OnTimeDB/part-00003
  inflating: /home/sergio_ibarra1795/OnTimeDB/part-00004
  inflating: /home/sergio_ibarra1795/OnTimeDB/part-00005
  inflating: /home/sergio_ibarra1795/OnTimeDB/part-00006
  inflating: /home/sergio_ibarra1795/OnTimeDB/_SUCCESS
total 3276
drwxrwxr-x 2 root          root          4096 Nov 23  2021 .
drwxr-xr-x 7 sergio_ibarra1795 sergio_ibarra1795 4096 Sep 23 18:24 ..
-rw-r--r-- 1 root          root          7428 Jul 18  2018
.part-00000.crc
-rw-r--r-- 1 root          root          2256 Jul 18  2018
.part-00001.crc
-rw-r--r-- 1 root          root          3184 Jul 18  2018
.part-00002.crc
-rw-r--r-- 1 root          root          2916 Jul 18  2018
.part-00003.crc
-rw-r--r-- 1 root          root          6948 Jul 18  2018
.part-00004.crc
-rw-r--r-- 1 root          root          2012 Jul 18  2018
.part-00005.crc
-rw-r--r-- 1 root          root          1076 Jul 18  2018
.part-00006.crc
-rw-r--r-- 1 root          root              0 Jul 18  2018 _SUCCESS
-rw-r--r-- 1 root          root      949404 Jul 18  2018 part-00000
-rw-r--r-- 1 root          root     287594 Jul 18  2018 part-00001
-rw-r--r-- 1 root          root     406415 Jul 18  2018 part-00002
-rw-r--r-- 1 root          root     372196 Jul 18  2018 part-00003
-rw-r--r-- 1 root          root     888266 Jul 18  2018 part-00004
-rw-r--r-- 1 root          root     256021 Jul 18  2018 part-00005
-rw-r--r-- 1 root          root     136582 Jul 18  2018 part-00006
mkdir: `/tmp/dcd/OnTimeDB': File exists
```

```
[7]: bd = sqlContext.read.csv("hdfs:///tmp/dcd/OnTimeDB/", inferSchema=True,
    ↪header=True)
sqlContext.registerDataFrameAsTable(bd, "bd")
bd.count()
```

[7]: 30466

```
[8]: bd.show(10)
```

```
[Stage 5:> (0 + 1) / 1]
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|Year|Month|DayofMonth|DayOfWeek|CRSDepTime|UniqueCarrier|TailNum|ArrDelay|DepDe
lay|Origin|Dest|Distance|Cancelled|Diverted|CarrierDelay|WeatherDelay|NASDelay|S
ecurityDelay|LateAircraftDelay|LogD|Retraso|RetrasoNeto|Horario|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|2016| 12| 1| 4| 845| AA| N8ARAA| -7.0|
-5.0| LAX| DFW| 1235.0| 0.0| 0.0| 0.0| 0.0| 0.0|
0.0| 0.0|3.0916669575956846| 0| -2.0| 2|
|2016| 12| 2| 5| 845| AA| N8ARAA| -3.0|
5.0| LAX| DFW| 1235.0| 0.0| 0.0| 0.0| 0.0| 0.0|
0.0| 0.0|3.0916669575956846| 0| -8.0| 2|
|2016| 12| 3| 6| 845| AA| N8ABAA| -3.0|
-3.0| LAX| DFW| 1235.0| 0.0| 0.0| 0.0| 0.0| 0.0|
0.0| 0.0|3.0916669575956846| 0| 0.0| 2|
|2016| 12| 4| 7| 845| AA| N8ABAA| -2.0|
-7.0| LAX| DFW| 1235.0| 0.0| 0.0| 0.0| 0.0| 0.0|
0.0| 0.0|3.0916669575956846| 0| 5.0| 2|
|2016| 12| 5| 1| 845| AA| N8ACAA| -2.0|
-6.0| LAX| DFW| 1235.0| 0.0| 0.0| 0.0| 0.0| 0.0|
0.0| 0.0|3.0916669575956846| 0| 4.0| 2|
|2016| 12| 6| 2| 845| AA| N867AA| 0.0|
-1.0| LAX| DFW| 1235.0| 0.0| 0.0| 0.0| 0.0| 0.0|
0.0| 0.0|3.0916669575956846| 0| 1.0| 2|
|2016| 12| 7| 3| 845| AA| N8ACAA| -6.0|
0.0| LAX| DFW| 1235.0| 0.0| 0.0| 0.0| 0.0| 0.0|
0.0| 0.0|3.0916669575956846| 0| -6.0| 2|
|2016| 12| 8| 4| 845| AA| N8AKAA| 7.0|
0.0| LAX| DFW| 1235.0| 0.0| 0.0| 0.0| 0.0| 0.0|
0.0| 0.0|3.0916669575956846| 0| 7.0| 2|
|2016| 12| 9| 5| 845| AA| N8ARAA| -9.0|
-1.0| LAX| DFW| 1235.0| 0.0| 0.0| 0.0| 0.0| 0.0|
0.0| 0.0|3.0916669575956846| 0| -8.0| 2|
```

2016	12	10	6	845	AA	N8AKAA	-2.0
-1.0	LAX	DFW	1235.0	0.0	0.0	0.0	0.0
0.0			0.0	3.0916669575956846	0	-1.0	2

```

+---+---+---+---+---+---+---+---+
--+---+---+---+---+---+---+---+
-----+---+---+---+---+---+---+

```

only showing top 10 rows

```
[9]: # Ver las variables disponibles
bd.dtypes
```

```
[9]: [('Year', 'int'),
      ('Month', 'int'),
      ('DayofMonth', 'int'),
      ('DayOfWeek', 'int'),
      ('CRSDepTime', 'int'),
      ('UniqueCarrier', 'string'),
      ('TailNum', 'string'),
      ('ArrDelay', 'double'),
      ('DepDelay', 'double'),
      ('Origin', 'string'),
      ('Dest', 'string'),
      ('Distance', 'double'),
      ('Cancelled', 'double'),
      ('Diverted', 'double'),
      ('CarrierDelay', 'double'),
      ('WeatherDelay', 'double'),
      ('NASDelay', 'double'),
      ('SecurityDelay', 'double'),
      ('LateAircraftDelay', 'double'),
      ('LogD', 'double'),
      ('Retraso', 'int'),
      ('RetrasoNeto', 'double'),
      ('Horario', 'int')]
```

```
[10]: spark.sql("SELECT year, Retraso, RetrasoNeto, Horario from bd LIMIT 10").show()
```

year	Retraso	RetrasoNeto	Horario
2016	0	-2.0	2
2016	0	-8.0	2
2016	0	0.0	2
2016	0	5.0	2
2016	0	4.0	2
2016	0	1.0	2

2016	0	-6.0	2
2016	0	7.0	2
2016	0	-8.0	2
2016	0	-1.0	2

+-----+-----+-----+

```
[11]: #Agregar una variable numerica (IndexUniqueCarrier) basada en la variable
      ↪ alfanumerica
      #de la compañía que opera el vuelo (UniqueCarrier)

      from pyspark.ml.feature import StringIndexer

      indexer =
        ↪ StringIndexer(inputCol='UniqueCarrier',outputCol='IndexUniqueCarrier') #el
        ↪ índice empieza en el 0!
      bd1=indexer.fit(bd).transform(bd)

      #Se muestra el numero de vuelos operados por cada compañía, la variable
      ↪ IndexUniqueCarrier
      #ya se puede utilizar en el modelo

      bd1.groupBy('UniqueCarrier','IndexUniqueCarrier').count().
        ↪ sort('IndexUniqueCarrier').show()
```

[Stage 12:=====>

(1 + 1) / 2]

UniqueCarrier	IndexUniqueCarrier	count
	AA	0.0 8853
	UA	1.0 6112
	WN	2.0 5395
	DL	3.0 4239
	VX	4.0 1703
	NK	5.0 1581
	F9	6.0 1295
	OO	7.0 1166
	B6	8.0 121
	EV	9.0 1

+-----+-----+-----+

1.1 Ajuste del modelo

```
[12]: from pyspark.ml.feature import VectorAssembler, StringIndexer
      from pyspark.sql.functions import col

      #Crear un arreglo de variables predictoras llamdo 'features'
      #ArrDelay representa el numero de minutos que un vuelo tiene de retraso
      #Si un vuelo llega con mas de 15 minutos de retraso la variable 'Retraso' tiene
      ↪valor 1
      #Renombrar la variable objetivo (Retraso) como 'label'

      #En el caso particular de los árboles de clasificación, la variable objetivo
      ↪debe ser de tipo doble.
      #Por lo tanto, transformar la variable a tipo doble.
      #Además, la variable debe estar convertida a través de la función stringIndexer
      #para poder ser analizada por el modelo
      #Por lo tanto, la variable de trabajo, en este caso, será 'label2'

      a1 = VectorAssembler(
          inputCols=['DepDelay', 'Distance', 'DayOfWeek',
                     'CRSDepTime', 'IndexUniqueCarrier'],
          outputCol='features')

      bd2 = a1.transform(bd1).select(col("Retraso").cast('double').
          ↪alias("label"), 'features')

      stringIndexer = StringIndexer(inputCol = 'label', outputCol = 'label2')
      sI = stringIndexer.fit(bd2)
      bd2 = sI.transform(bd2)
      bd2.dtypes
```

```
[12]: [('label', 'double'), ('features', 'vector'), ('label2', 'double')]
```

```
[14]: #Mostrar un solo renglon de la BD
      bd2.show(5)
```

```
[Stage 19:> (0 + 1) / 1]
```

```
+-----+-----+-----+
|label|          features|label2|
+-----+-----+-----+
| 0.0|[-5.0,1235.0,4.0,...| 0.0|
| 0.0|[5.0,1235.0,5.0,8...| 0.0|
| 0.0|[-3.0,1235.0,6.0,...| 0.0|
| 0.0|[-7.0,1235.0,7.0,...| 0.0|
| 0.0|[-6.0,1235.0,1.0,...| 0.0|
```

+-----+-----+-----+
only showing top 5 rows

1.1.1 Partición Test - Train

```
[15]: #70% Train
      #30% Test
      (bd_train, bd_test) = bd2.randomSplit([0.7, 0.3],seed=123)
      print("Renglones de la BD Train: ", bd_train.count())
      print("Renglones de la BD Test: ",bd_test.count())
```

Renglones de la BD Train: 21219

[Stage 23:=====> (1 + 1) / 2]

Renglones de la BD Test: 9247

```
[16]: # Utilizamos el modelo DecisionTreeClassifier para generar un prediccion basado
      ↪ en los features
      # disponibles
      #Con una profundidad (maxDepth) de 5
      #Especificar la variable objetivo (labelCol)

      from pyspark.ml.classification import DecisionTreeClassifier as DTC

      rt = DTC(maxDepth=5, labelCol = 'label12')

      model = rt.fit(bd_train)
      pred = model.transform(bd_train)
```

```
[17]: #La columna rawPrediction, está especificando el número de casos negativos y
      ↪ positivos
      #en cada uno de los nodos terminales pertinentes para cada observación.
      #[Casos para 0, Casos para 1]

      #El campo probability muestra la probabilidad de ser 0 o 1
      #El valor predicho, estableciendo un punto de corte del 50% se muestra en el
      ↪ campo prediction

      pred.show()
```

+-----+-----+-----+-----+-----+-----+-----+
--+

```

label|          features|label2|  rawPrediction|
probability|prediction|
+-----+-----+-----+-----+-----+-----+
--+
|  0.0| [-21.0,868.0,6.0,...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-20.0,1440.0,6.0...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-19.0,1440.0,3.0...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-18.0,602.0,5.0,...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-17.0,888.0,6.0,...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-17.0,1440.0,1.0...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-17.0,1744.0,1.0...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-16.0,641.0,6.0,...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-16.0,868.0,6.0,...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-15.0,731.0,1.0,...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-15.0,868.0,3.0,...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-15.0,888.0,4.0,...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-15.0,888.0,5.0,...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-15.0,1464.0,6.0...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-15.0,1514.0,2.0...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-15.0,1514.0,4.0...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-14.0,236.0,2.0,...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-14.0,236.0,3.0,...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-14.0,236.0,4.0,...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
|  0.0| [-14.0,255.0,3.0,...|    0.0| [13665.0,985.0] | [0.93276450511945...|
0.0|
+-----+-----+-----+-----+-----+-----+
--+
only showing top 20 rows

```



```
[18]: #Validar el modelo
      #Calcular el Area bajo la curva
      #El AUC proporciona una medición agregada del rendimiento en todos los umbrales
      ↪ de clasificación
      #posibles

      from pyspark.ml.evaluation import BinaryClassificationEvaluator as BCE
      print('AUC=',BCE(metricName="areaUnderROC", rawPredictionCol = 'probability').
      ↪ evaluate(pred))
```

AUC= 0.8966334172544019

```
[19]: #Generar una tabla de frecuencias de las distintas probabilidades, es decir de
      ↪ los distintos
      #nodos terminales

      pred.groupBy('probability').count().sort('count').show(50)
```

[Stage 52:=====> (1 + 1) / 2]

```
+-----+-----+
|      probability|count|
+-----+-----+
|      [1.0,0.0]|    2|
| [0.65116279069767...|   43|
| [0.35526315789473...|   76|
| [0.25984251968503...|  127|
| [0.66666666666666...|  144|
| [0.50531914893617...|  188|
| [0.09150326797385...|  306|
| [0.46683673469387...|  392|
| [0.27331887201735...|  461|
| [0.70650032829940...| 1523|
| [0.01118838826731...| 3307|
| [0.93276450511945...|14650|
+-----+-----+
```

```
[20]: #Generar la matriz de confusion
      pred.groupBy('label','prediction').count().show()
```

[Stage 55:=====> (1 + 1) / 2]

```
+-----+-----+
|label|prediction|count|
```

```

+-----+-----+-----+
| 1.0|      1.0| 4235|
| 0.0|      1.0|  434|
| 1.0|      0.0| 1588|
| 0.0|      0.0|14962|
+-----+-----+-----+

```

```

[21]: #Generar algunas estadísticas para tener una idea de cómo fueron las
      ↳predicciones

numSuccesses = pred.where("""(prediction = 0.0 AND label2 = 0.0) OR (prediction
      ↳= 1.0 AND label2 = 1.0)""").count()
numInspections = pred.count()

print ("Se realizaron", numInspections, "inspecciones y existen", numSuccesses,
      ↳"predicciones exitosas")
print ("Esta es una tasa de éxito del", str((float(numSuccesses) /
      ↳float(numInspections)) * 100) + "%")

```

[Stage 61:=====> (1 + 1) / 2]

Se realizaron 21219 inspecciones y existen 19197 predicciones exitosas
 Esta es una tasa de éxito del 90.47080446769404%

```

[22]: # DecisionTreeClassifier(featuresCol="features",
      #   labelCol="label",
      #   predictionCol="prediction",
      #   probabilityCol="probability",
      #   rawPredictionCol="rawPrediction",
      #   maxDepth=5,
      #   maxBins=32,
      #   minInstancesPerNode=1,
      #   minInfoGain=0.0,
      #   maxMemoryInMB=256,
      #   impurity="gini" / impurity="entropy" )

```

```

[23]: #Cambiando las propiedades del ejecutamos un nuevo arbol con una profundidad de
      ↳20

rt = DTC(maxDepth=20, labelCol = 'label2')
model = rt.fit(bd_train)
pred = model.transform(bd_train)

#Evaluar el modelo con AUC que ahora ha aumentado

```

```
print('AUC=',BCE(metricName="areaUnderROC", rawPredictionCol = 'probability').  
      ↪evaluate(pred))
```

AUC= 0.9990185263935085

1.1.2 Validación externa

```
[24]: #Validación externa con la BD test  
  
predtest = model.transform(bd_test)  
  
print('AUC=',BCE(metricName="areaUnderROC",rawPredictionCol = 'probability').  
      ↪evaluate(predtest))
```

AUC= 0.8193169919620452

[]: